

Micah Johnson

Kennedale, TX · micah@createsomething.xyz · linkedin.com/in/micahryanjohnson

January 30, 2026

Reveal Technology

Engineering Team

I'm applying for the AI Engineer role at Reveal. Your concept of "earned delegation"—AI systems that start as advisors, mature into copilots, and ultimately become trusted operators—resonates deeply with how I think about automation. The goal isn't "AI assistance." It's progressive trust that's bounded, auditable, and aligned.

On multi-step agentic systems: I'm currently building WORKWAY, an agentic workflow platform. I've implemented orchestration patterns for sequential workflows (Harness), parallel execution (Gastown), and iterative refinement (Ralph). These aren't theoretical—they're production systems with state management, rollback capabilities, and audit trails. I understand the difference between a demo that works and a system you can actually trust with real-world consequences.

On earned delegation and trust modeling: My career is literally an example of this concept. I started in customer support at Webflow—handling 500+ direct user interactions, learning what trust looks like from the user's perspective. I built an internal LMS, proved competence, got promoted to team manager. The Marketplace team noticed my work, asked me to build their systems. I left, got brought back as a contractor specifically for that work, then rejoined full-time. Each step was earned through demonstrated competence. I understand why progressive autonomy matters because I've lived it.

On iterating based on behavior, not theory: Two years in support taught me that users don't behave the way you expect. You can't design trust models from first principles—you have to observe how people actually interact with systems and adapt. When you list "willingness to engage deeply with users" as a preference, that's not optional for me. It's how I've always worked.

On the technical stack: Python daily, LLM orchestration and RAG pipelines in production, tool calling and function execution, Docker and CI/CD automation. I've optimized for latency, cost, and reliability in real systems. The defense context is new, but the core patterns—constraint-aware reasoning, audibility, explainability—are exactly what I've been building.

Reveal's mission—intelligent automation that replaces manual work while remaining transparent and aligned—is the right framing for AI deployment. I'd be excited to help build agentic systems that earn trust through demonstrated competence, not just claimed capability.

Micah Johnson